

```
1 from google.colab import files
2 uploaded = files.upload()
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to

Saving dummydata.csv to dummydata.csv

```
1 import pandas as pd
2 import numpy as np
3 df = pd.read_csv('dummydata.csv')
4 pd.DataFrame.from_records(df)
5 df.head()
```

	Unnamed: 0	age	self_employed	family_history	mh_treatment	interfere	company_size	remote	tech_company	mh_negati
0	1	3	1	1	0	3	1	1	1	
1	2	2	0	1	1	3	4	0	1	
2	3	2	1	0	0	1	1	1	1	
3	4	3	0	0	1	4	3	1	1	
4	5	3	0	0	1	1	6	0	0	

```
1 df.describe
```

<bound method NDFrame.describe of

	Unnamed: 0	age	...	country_Ireland	country_India
0	1	3	...	0	0
1	2	2	...	0	0
2	3	2	...	0	0
3	4	3	...	0	0
4	5	3	...	0	0
..	...	...	...	...	...
891	892	2	...	0	0
892	893	2	...	0	0
893	894	2	...	0	0
894	895	2	...	0	0
895	896	3	...	0	0

[896 rows x 51 columns]>

```
1 df.columns
```

Index(['Unnamed: 0', 'age', 'self\_employed', 'family\_history', 'mh\_treatment', 'interfere', 'company\_size', 'remote', 'tech\_company', 'mh\_negative\_consequence\_flag', 'ph\_negative\_consequence\_flag', 'mh\_disscuss\_coworker', 'mh\_disscuss\_supervisor', 'interview\_mh\_bringup', 'interview\_ph\_bringup', 'witness\_mh\_nc', 'anonymity\_protected\_Yes', 'anonymity\_protected\_No', 'anonymity\_protected\_Don't know', 'awareness\_mh\_benefits\_Not sure', 'awareness\_mh\_benefits\_Yes', 'awareness\_mh\_benefits\_No', 'gender\_M', 'gender\_F', 'gender\_T', 'medical\_leave\_easy\_Very easy', 'medical\_leave\_easy\_Somewhat difficult', 'medical\_leave\_easy\_Don't know', 'medical\_leave\_easy\_Very difficult', 'medical\_leave\_easy\_Somewhat easy', 'mh\_benefits\_Yes', 'mh\_benefits\_No', 'mh\_benefits\_Don't know', 'mh\_discuss\_Yes', 'mh\_discuss\_No', 'mh\_discuss\_Don't know', 'mh\_resources\_Don't know', 'mh\_resources\_No', 'mh\_resources\_Yes', 'mh\_serious\_ph\_Yes', 'mh\_serious\_ph\_No', 'mh\_serious\_ph\_Don't know', 'country\_United States', 'country\_United Kingdom', 'country\_Canada', 'country\_Netherlands', 'country\_Australia', 'country\_France', 'country\_Germany', 'country\_Ireland', 'country\_India'], dtype='object')

```
1 df.drop(df.columns[[0]], axis=1, inplace=True)
2 df.head()
```

	age	self_employed	family_history	mh_treatment	interfere	company_size	remote	tech_company	mh_negative_consequ
0	3	1	1	0	3	1	1	1	
1	2	0	1	1	3	4	0	1	
2	2	1	0	0	1	1	1	1	
3	3	0	0	1	4	3	1	1	
4	3	0	0	1	1	6	0	0	

```
1 df.shape
```

(896, 50)

Supervised ML

```
1 X=df.loc[:, df.columns != 'mh_treatment']
2 X.head()
```

	age	self_employed	family_history	interfere	company_size	remote	tech_company	mh_negative_consequence_flag	ph_
0	3	1	1	3	1	1	1	0	
1	2	0	1	3	4	0	1	1	
2	2	1	0	1	1	1	1	0	
3	3	0	0	4	3	1	1	1	
4	3	0	0	1	6	0	0	1	

```
1 X.columns
```

Index(['age', 'self\_employed', 'family\_history', 'interfere', 'company\_size', 'remote', 'tech\_company', 'mh\_negative\_consequence\_flag', 'ph\_negative\_consequence\_flag', 'mh\_disscuss\_coworker', 'mh\_disscuss\_supervisor', 'interview\_mh\_bringup', 'interview\_ph\_bringup', 'witness\_mh\_nc', 'anonymity\_protected\_Yes', 'anonymity\_protected\_No', 'anonymity\_protected\_Don't know', 'awareness\_mh\_benefits\_Not sure', 'awareness\_mh\_benefits\_Yes', 'awareness\_mh\_benefits\_No', 'gender\_M', 'gender\_F', 'gender\_T', 'medical\_leave\_easy\_Very easy', 'medical\_leave\_easy\_Somewhat difficult', 'medical\_leave\_easy\_Don't know', 'medical\_leave\_easy\_Very difficult', 'medical\_leave\_easy\_Somewhat easy', 'mh\_benefits\_Yes', 'mh\_benefits\_No', 'mh\_benefits\_Don't know', 'mh\_discuss\_Yes', 'mh\_discuss\_No', 'mh\_discuss\_Don't know', 'mh\_resources\_Don't know', 'mh\_resources\_No', 'mh\_resources\_Yes', 'mh\_serious\_ph\_Yes', 'mh\_serious\_ph\_No', 'mh\_serious\_ph\_Don't know', 'country\_United States', 'country\_United Kingdom', 'country\_Canada', 'country\_Netherlands', 'country\_Australia', 'country\_France', 'country\_Germany', 'country\_Ireland', 'country\_India'], dtype='object')

```
1 y=df.mh_treatment
2 y.head()
```

0 0
1 1
2 0
3 1
4 1
Name: mh\_treatment, dtype: int64

```
1 np.random.seed(888)
2 #Split X and y into train (70%) and test (30%) sets.
3 from sklearn.model_selection import train_test_split
4 # Let 20% of the data to be a test set
5 train_x, test_x, train_y, test_y = train_test_split(X, y, test_size=0.30, random_state=0)
```

1. Linear Regression

```
1 #fit the model
2 from sklearn.linear_model import LinearRegression
3 from sklearn.metrics import mean_squared_error
4 modell = LinearRegression()
5 modell.fit(train_x,train_y)
```

LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=None, normalize=False)

```
1 # coefficeints of the trained model
2 print('\nCoefficient of model :', modell.coef_)
3
4 # intercept of the model
5 print('\nIntercept of model',modell.intercept_)
6
7 # predict thetarget on the test dataset
```

```

8 predict_train1 = model1.predict(train_x)
9
10 # Root Mean Squared Error on training dataset
11 rmse_train1 = mean_squared_error(train_y,predict_train1)**(0.5)
12 print('\nRMSE on train dataset : ', rmse_train1)
13

```



Coefficient of model : [ 7.94394178e-02 -3.13073938e-02 1.35247094e-01 2.00005298e-01  
 2.16494895e-03 6.28433680e-03 -3.96517007e-03 2.29191247e-02  
 -9.32070297e-03 6.02567060e-02 -8.94607436e-03 9.01932694e-03  
 1.84655909e-02 6.08269687e-03 2.47073445e+11 2.47073445e+11  
 2.47073445e+11 -6.17962248e+11 -6.17962248e+11 -6.17962248e+11  
 -2.39548373e+11 -2.39548373e+11 -2.39548373e+11 -3.63133064e+11  
 -3.63133064e+11 -3.63133064e+11 -3.63133064e+11 -3.63133064e+11  
 -4.70124964e+11 -4.70124964e+11 -4.70124964e+11 -5.22730142e+11  
 -5.22730142e+11 -5.22730142e+11 -7.48774855e+11 -7.48774855e+11  
 -7.48774855e+11 -1.77157781e+11 -1.77157781e+11 -1.77157781e+11  
 -1.33097841e+11 -1.33097841e+11 -1.33097841e+11 -1.33097841e+11  
 -1.33097841e+11 -1.33097841e+11 -1.33097841e+11 -1.33097841e+11  
 -1.33097841e+11]

Intercept of model 3025455823032.587

RMSE on train dataset : 0.38898088972415895

```

1 # predict the target on the testing dataset
2 predict_test1 = model1.predict(test_x)
3
4 # Root Mean Squared Error on testing dataset
5 rmse_test1 = mean_squared_error(test_y,predict_test1)**(0.5)
6 print('\nRMSE on test dataset : ', rmse_test1)

```



RMSE on test dataset : 0.3846918650527029

## 2. Tree

```

1 from sklearn.tree import DecisionTreeClassifier
2 from sklearn.metrics import accuracy_score
3 model2 = DecisionTreeClassifier()
4
5 # fit the model with the training data
6 model2.fit(train_x,train_y)
7
8 # depth of the decision tree
9 print('Depth of the Decision Tree :', model2.get_depth())
10

```



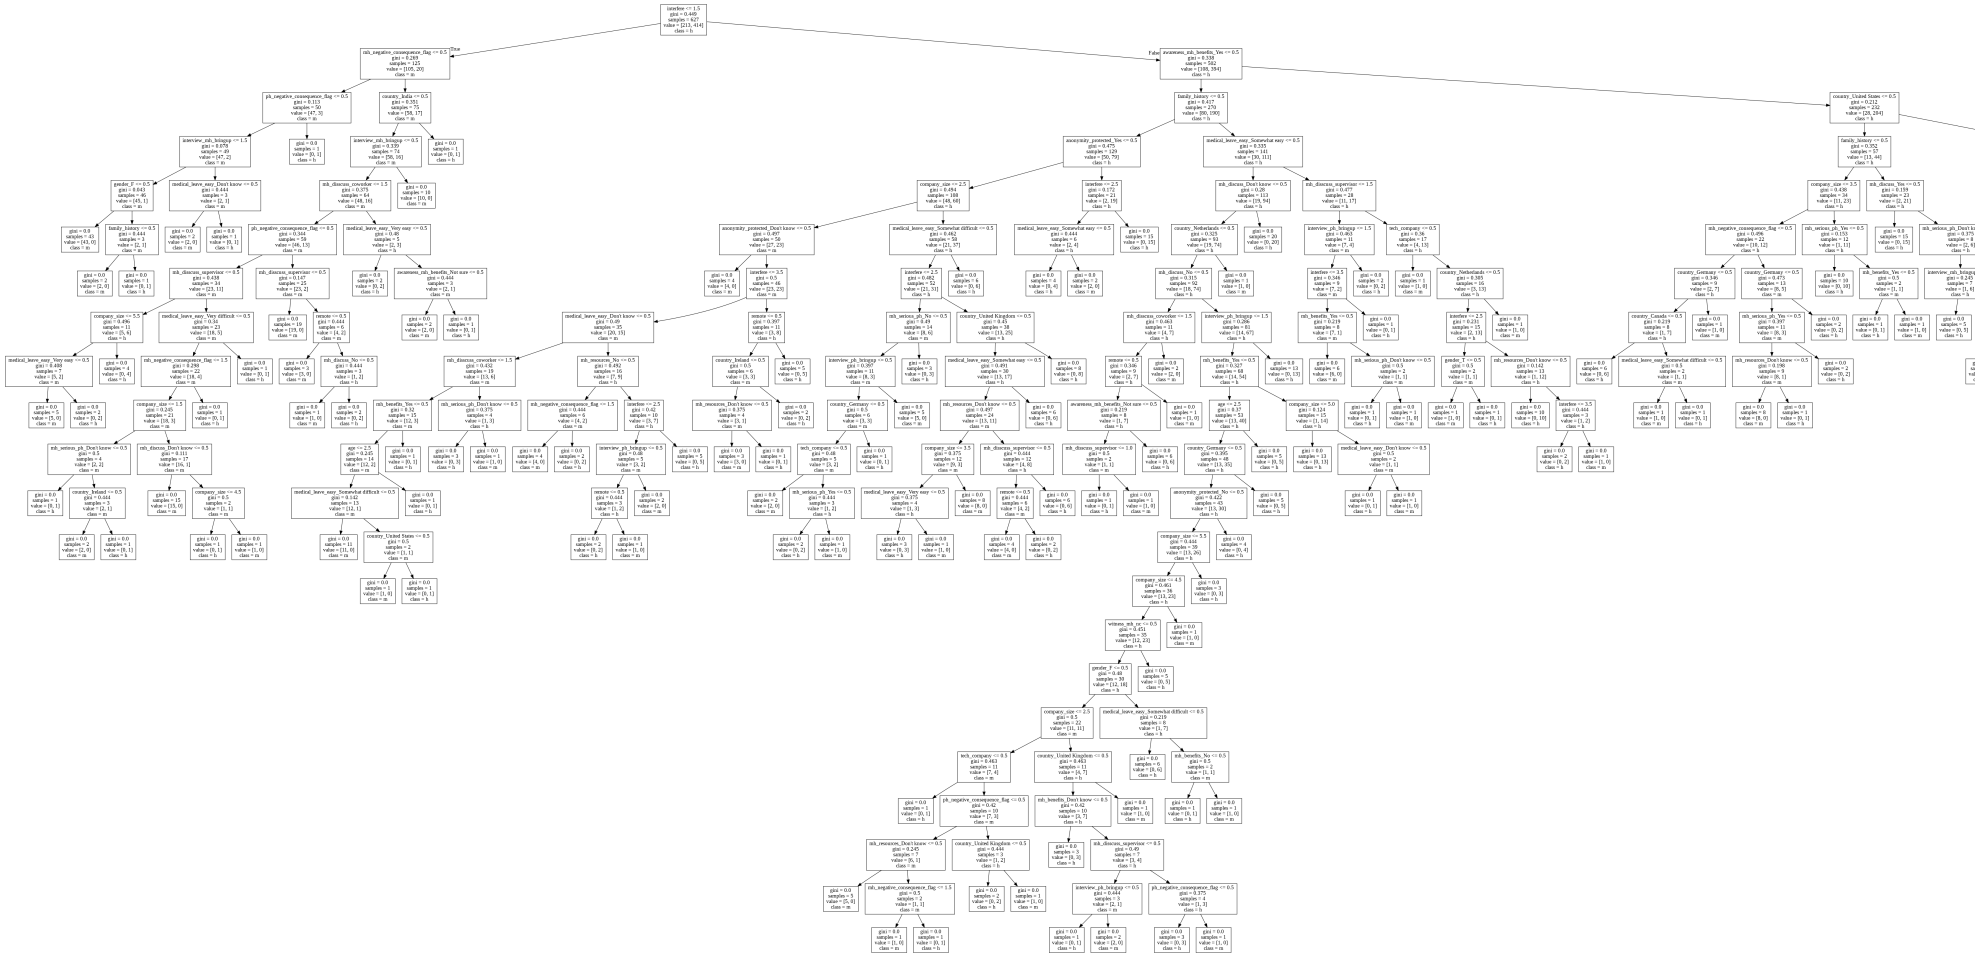
Depth of the Decision Tree : 21

```


1 from sklearn import datasets
2 from IPython.display import Image
3 from sklearn import tree
4 import pydotplus
5 # Create DOT data
6 dot_data = tree.export_graphviz(model2.fit(train_x,train_y), out_file=None,
7                                 feature_names=train_x.columns,
8                                 class_names=train_y.name)
9
10 # Draw graph
11 graph = pydotplus.graph_from_dot_data(dot_data)
12
13 # Show graph
14 Image(graph.create_png())

```





```
1 # predict the target on the train dataset
2 predict_train2 = model2.predict(train_x)
3
4 # Accuracy Score on train dataset
5 accuracy_train2 = accuracy_score(train_y,predict_train2)
6 print('accuracy_score on train dataset : ', accuracy_train2)
7
```

 accuracy\_score on train dataset : 1.0

```
1 # predict the target on the test dataset
2 predict_test2 = model2.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test2 = accuracy_score(test_y,predict_test2)
6 print('accuracy_score on test dataset : ', accuracy_test2)
```

 accuracy\_score on test dataset : 0.7026022304832714

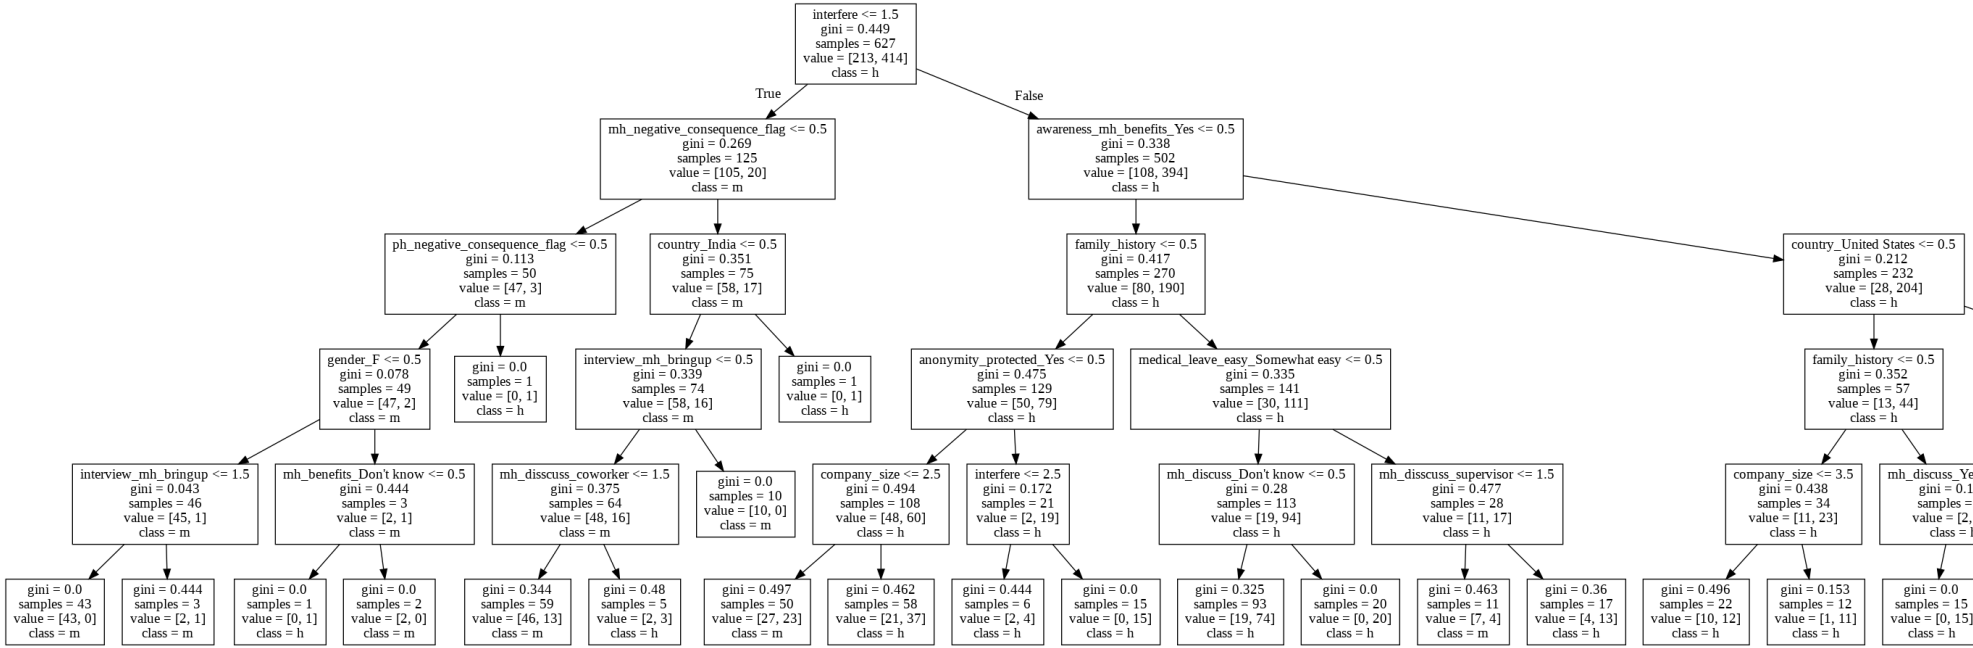
3. Decision Tree Calssifier

```
1 model3 = DecisionTreeClassifier(max_depth=5)
2
3 # fit the model with the training data
4 model3.fit(train_x,train_y)
5
6 # depth of the decision tree
7 print('Depth of the Decision Tree :', model3.get_depth())
```

 Depth of the Decision Tree : 5

```
1 # Create DOT data
2 dot_data3 = tree.export_graphviz(model3.fit(train_x,train_y), out_file=None,
3                                 feature_names=train_x.columns,
4                                 class_names=train_y.name)
5
6 # Draw graph
7 graph3 = pydotplus.graph_from_dot_data(dot_data3)
8
9 # Show graph
10 Image(graph3.create_png())
```





```
1 # predict the target on the train dataset
2 predict_train3 = model3.predict(train_x)
3
4 # Accuray Score on train dataset
5 accuracy_train3 = accuracy_score(train_y,predict_train3)
6 print('accuracy_score on train dataset : ', accuracy_train3)
```

accuracy\_score on train dataset : 0.8133971291866029

```
1 # predict the target on the test dataset
2 predict_test3 = model3.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test3 = accuracy_score(test_y,predict_test3)
6 print('accuracy_score on test dataset : ', accuracy_test3)
```

accuracy\_score on test dataset : 0.758364312267658

4. SVM

```
1 from sklearn.svm import SVC
2 model4 = SVC()
3
4 # fit the model with the training data
5 model4.fit(train_x,train_y)
```

SVC(C=1.0, break\_ties=False, cache\_size=200, class\_weight=None, coef0=0.0, decision\_function\_shape='ovr', degree=3, gamma='scale', kernel='rbf', max\_iter=-1, probability=False, random\_state=None, shrinking=True, tol=0.001, verbose=False)

```
1 # predict the target on the train dataset
2 predict_train4 = model4.predict(train_x)
3
4 # Accuray Score on train dataset
5 accuracy_train4 = accuracy_score(train_y,predict_train4)
6 print('accuracy_score on train dataset : ', accuracy_train4)
```

accuracy\_score on train dataset : 0.8181818181818182

```
1 # predict the target on the test dataset
2 predict_test4 = model4.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test4 = accuracy_score(test_y,predict_test4)
6 print('accuracy_score on test dataset : ', accuracy_test4)
```

accuracy\_score on test dataset : 0.8066914498141264

5. KNN

```
1 from sklearn.neighbors import KNeighborsClassifier
2 model5 = KNeighborsClassifier()
3
4 # fit the model with the training data
5 model5.fit(train_x,train_y)
```

```

5 model5.fit(train_x, train_y)
6
7 # Number of Neighbors used to predict the target
8 print('\nThe number of neighbors used to predict the target : ',model5.n_neighbors)

```



The number of neighbors used to predict the target : 5

```

1 # predict the target on the train dataset
2 predict_train5 = model5.predict(train_x)
3
4 # Accuray Score on train dataset
5 accuracy_train5 = accuracy_score(train_y,predict_train5)
6 print('accuracy_score on train dataset : ', accuracy_train5)

```



accuracy\_score on train dataset : 0.8197767145135566

```

1 # predict the target on the test dataset
2 predict_test5 = model5.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test5 = accuracy_score(test_y,predict_test5)
6 print('accuracy_score on test dataset : ', accuracy_test5)

```



accuracy\_score on test dataset : 0.7211895910780669

## 6. GBM

```

1 from sklearn.ensemble import GradientBoostingClassifier
2 model6 = GradientBoostingClassifier(n_estimators=100,max_depth=5)
3
4 # fit the model with the training data
5 model6.fit(train_x,train_y)

```



```

GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=5,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=None, subsample=1.0, tol=0.0001,
                           validation_fraction=0.1, verbose=0,
                           warm_start=False)

```

```

1 # predict the target on the train dataset
2 predict_train6 = model6.predict(train_x)
3
4 # Accuray Score on train dataset
5 accuracy_train6 = accuracy_score(train_y,predict_train6)
6 print('\naccuracy_score on train dataset : ', accuracy_train6)

```



accuracy\_score on train dataset : 0.9904306220095693

```

1 # predict the target on the test dataset
2 predict_test6 = model6.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test6 = accuracy_score(test_y,predict_test6)
6 print('\naccuracy_score on test dataset : ', accuracy_test6)

```



accuracy\_score on test dataset : 0.724907063197026

## 7. XGBoost

```

1 from xgboost import XGBClassifier
2 model7 = XGBClassifier()
3
4 # fit the model with the training data
5 model7.fit(train_x,train_y)

```





```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0,
              learning_rate=0.1, max_delta_step=0, max_depth=3,
              min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
              nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=None, subsample=1, verbosity=1)
```

```
1 # predict the target on the train dataset
2 predict_train7 = model7.predict(train_x)
3
4 # Accuray Score on train dataset
5 accuracy_train7 = accuracy_score(train_y,predict_train7)
6 print('\naccuracy_score on train dataset : ', accuracy_train7)
```



```
accuracy_score on train dataset :  0.8389154704944178
```

```
1 # predict the target on the test dataset
2 predict_test7 = model7.predict(test_x)
3
4 # Accuracy Score on test dataset
5 accuracy_test7 = accuracy_score(test_y,predict_test7)
6 print('\naccuracy_score on test dataset : ', accuracy_test7)
```



```
accuracy_score on test dataset :  0.7695167286245354
```


## Unsupervised ML

### PCA + kmeans

```
1 #scaler data
2 from sklearn.preprocessing import MinMaxScaler
3 scaler = MinMaxScaler()
4 data_rescaled = scaler.fit_transform(df)
```

```
1 # find the suitable number of components
2 from sklearn.decomposition import PCA
3 pca = PCA().fit(data_rescaled)
4
5 % matplotlib inline
6 import matplotlib.pyplot as plt
7 plt.rcParams["figure.figsize"] = (12,6)
8
9 fig, ax = plt.subplots()
10 xi = np.arange(1, 21, step=1)
11 y = np.cumsum(pca.explained_variance_ratio_)[1:21]
12
13 plt.ylim(0.0,1.1)
14 plt.plot(xi, y, marker='o', linestyle='--', color='b')
15
16 plt.xlabel('Number of Components')
17 plt.xticks(np.arange(0, 21, step=1))
18 plt.ylabel('Cumulative variance (%)')
19 plt.title('The number of components needed to explain variance')
20
21 plt.axhline(y=0.70, color='r', linestyle='-')
22 plt.text(0.5, 0.85, '70% cut-off threshold', color = 'red', fontsize=16)
23
24 ax.grid(axis='x')
25 plt.show()
```





In order to get around 70%, we decided to choose the first 12 components.

```
1 #fit the model
2 pca_final = PCA(n_components = 12)
3 pca_final.fit(data_rescaled)
4 data_pca = pca_final.fit_transform(data_rescaled)
5 data_pca.shape
```


(896, 12)

```
1 #find reasonable number of k with elbow method
2 from sklearn.cluster import KMeans
3 from sklearn import metrics
4 from scipy.spatial.distance import cdist
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 distortions = []
9 inertias = []
10 mapping1 = {}
11 mapping2 = {}
12 K = range(1,10)
13
14 for k in K:
15     #Building and fitting the model
16     kmeanModel = KMeans(n_clusters=k).fit(data_pca)
17     kmeanModel.fit(data_pca)
18
19     distortions.append(sum(np.min(cdist(data_pca, kmeanModel.cluster_centers_,
20                                     'euclidean'),axis=1)) / data_pca.shape[0])
21     inertias.append(kmeanModel.inertia_)
22
23     mapping1[k] = sum(np.min(cdist(data_pca, kmeanModel.cluster_centers_,
24                                     'euclidean'),axis=1)) / data_pca.shape[0]
25     mapping2[k] = kmeanModel.inertia_
26
```

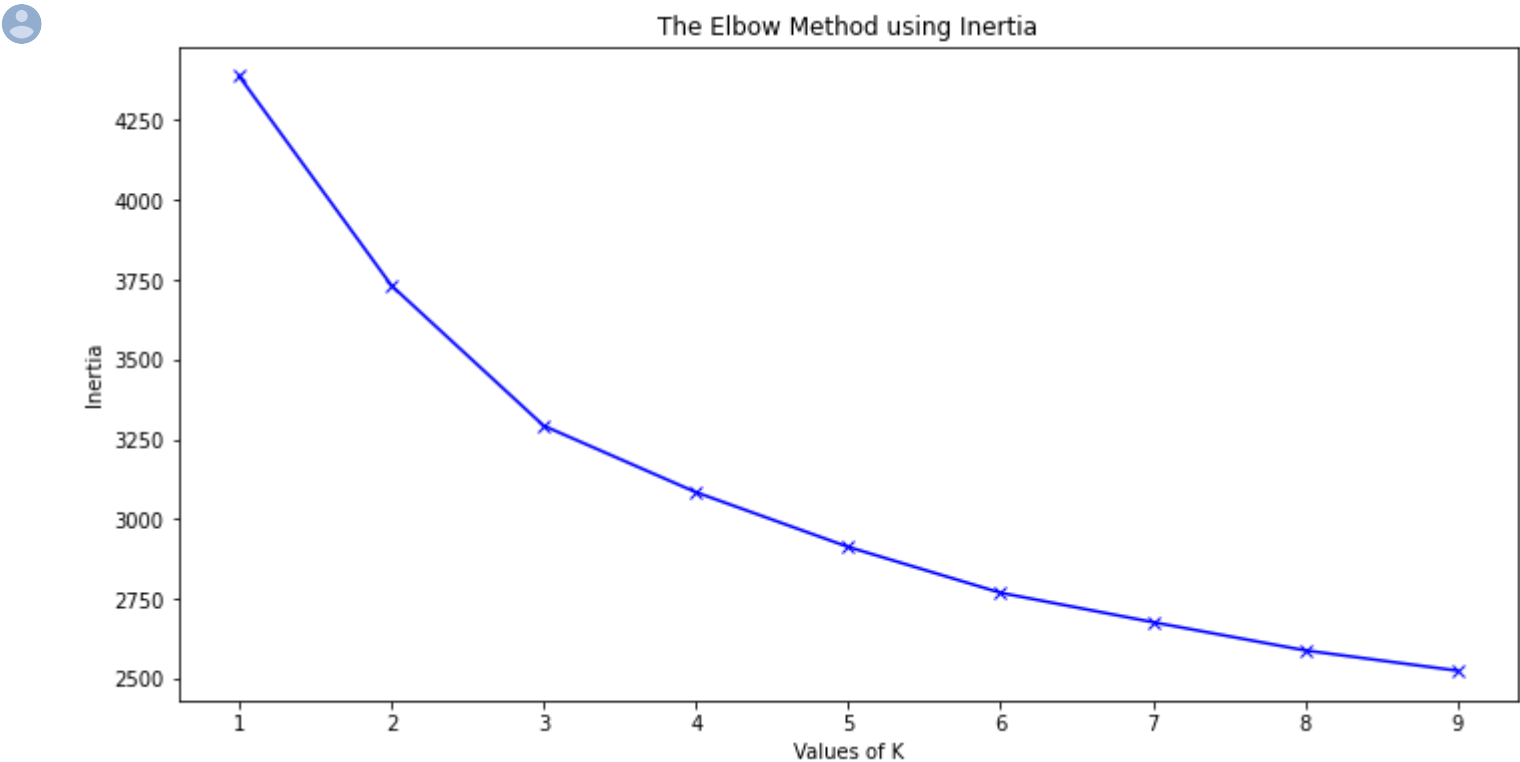
```
1 plt.plot(K, distortions, 'bx-')
2 plt.xlabel('Values of K')
3 plt.ylabel('Distortion')
4 plt.title('The Elbow Method using Distortion')
5 plt.show()
```







```
1 plt.plot(K, inertias, 'bx-')
2 plt.xlabel('Values of K')
3 plt.ylabel('Inertia')
4 plt.title('The Elbow Method using Inertia')
5 plt.show()
```



Based on the results, we chose to use k=5.


```
1 #fit the kmeans model
2 kmeans= KMeans(n_clusters=5)
3 kmeans5=kmeans.fit_predict(data_pca)
```

```
1 kmeans.cluster_centers_
```

```
array([[ -0.4429739 ,  1.05666511,  0.43913159,  0.28931182,  0.11870278,
        -0.09887435, -0.00617988, -0.07980882, -0.07975342, -0.05690273,
         0.12475788,  0.15104256],
       [ -1.07536002, -0.34658112, -0.08408743, -0.11913515,  0.10029392,
         0.18634079, -0.02428671,  0.00897236,  0.06673901,  0.05804446,
        -0.03544    , -0.17652532],
       [  1.53302396,  0.07519759,  0.0984913 , -0.14740826,  0.2511941 ,
         0.16342133, -0.03881158, -0.01913487, -0.05577262,  0.09761498,
         0.04129115, -0.1241347 ],
       [ -0.00363126, -1.22165991,  0.41078203,  0.09182208, -0.23131532,
        -0.18612103,  0.15502751, -0.00565208, -0.06049708, -0.10816998,
         0.06534425,  0.0312912 ],
       [  0.26556962,  0.15832695, -0.79598708, -0.08762594, -0.33144407,
        -0.14719943, -0.04429081,  0.0998988 ,  0.10781135, -0.02589955,
        -0.18514208,  0.15876326]])
```

```
1 #visualizing clustering
2 plt.scatter(data_pca[:,0],data_pca[:,1],c=kmeans5,cmap="rainbow")
```

<matplotlib.collections.PathCollection at 0x7f457c7b1198>



```
1 # Get cluster assignment labels
2 labels = kmeans.labels_
3 df['Group']=labels
4 df.head()
```

	age	self_employed	family_history	mh_treatment	interfere	company_size	remote	tech_company	mh_negative_consequ
0	3	1	1	0	3	1	1	1	
1	2	0	1	1	3	4	0	1	
2	2	1	0	0	1	1	1	1	
3	3	0	0	1	4	3	1	1	
4	3	0	0	1	1	6	0	0	

```
1 df.groupby('Group').describe()
```

	age								self_employed								family_history			
	count	mean	std	min	25%	50%	75%	max	count	mean	std	min	25%	50%	75%	max	count	mean		
Group																				
0	190.0	2.100000	0.378664	1.0	2.0	2.0	2.0	3.0	190.0	0.100000	0.300793	0.0	0.0	0.0	0.0	1.0	190.0	0.384211		
1	218.0	2.059633	0.347637	1.0	2.0	2.0	2.0	3.0	218.0	0.114679	0.319367	0.0	0.0	0.0	0.0	1.0	218.0	0.371560		
2	178.0	2.207865	0.433805	2.0	2.0	2.0	2.0	4.0	178.0	0.061798	0.241467	0.0	0.0	0.0	0.0	1.0	178.0	0.500000		
3	136.0	2.080882	0.323297	1.0	2.0	2.0	2.0	3.0	136.0	0.367647	0.483947	0.0	0.0	0.0	1.0	1.0	136.0	0.433824		
4	174.0	2.132184	0.339668	2.0	2.0	2.0	2.0	3.0	174.0	0.011494	0.106901	0.0	0.0	0.0	0.0	1.0	174.0	0.637931		

5 rows x 400 columns

```
1 df.groupby('Group').mean()
```

	age	self_employed	family_history	mh_treatment	interfere	company_size	remote	tech_company	mh_negativ
Group									
0	2.100000	0.100000	0.384211	0.505263	2.257895	3.573684	0.315789	0.821053	
1	2.059633	0.114679	0.371560	0.481651	2.481651	2.834862	0.288991	0.834862	
2	2.207865	0.061798	0.500000	0.735955	2.477528	4.516854	0.224719	0.707865	
3	2.080882	0.367647	0.433824	0.727941	2.801471	2.066176	0.514706	0.897059	
4	2.132184	0.011494	0.637931	0.890805	2.816092	4.097701	0.229885	0.816092	

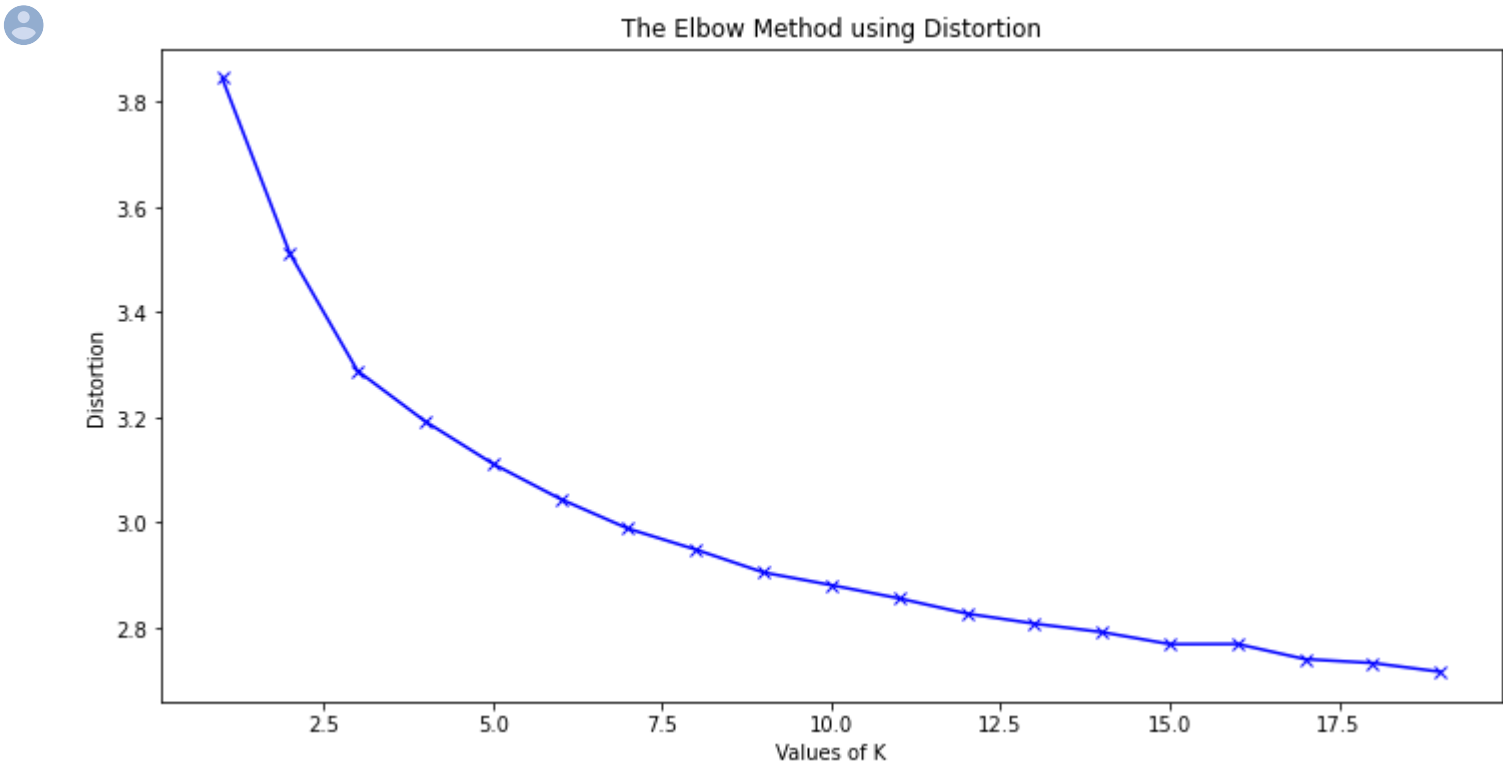
```
1 df.groupby('Group').mean().rank()
```

	age	self_employed	family_history	mh_treatment	interfere	company_size	remote	tech_company	mh_negative_con
Group									
0	3.0	3.0	2.0	2.0	1.0	3.0	4.0	3.0	
1	1.0	4.0	1.0	1.0	3.0	2.0	3.0	4.0	
2	5.0	2.0	4.0	4.0	2.0	5.0	1.0	1.0	
3	2.0	5.0	3.0	3.0	4.0	1.0	5.0	5.0	
4	4.0	1.0	5.0	5.0	5.0	4.0	2.0	2.0	

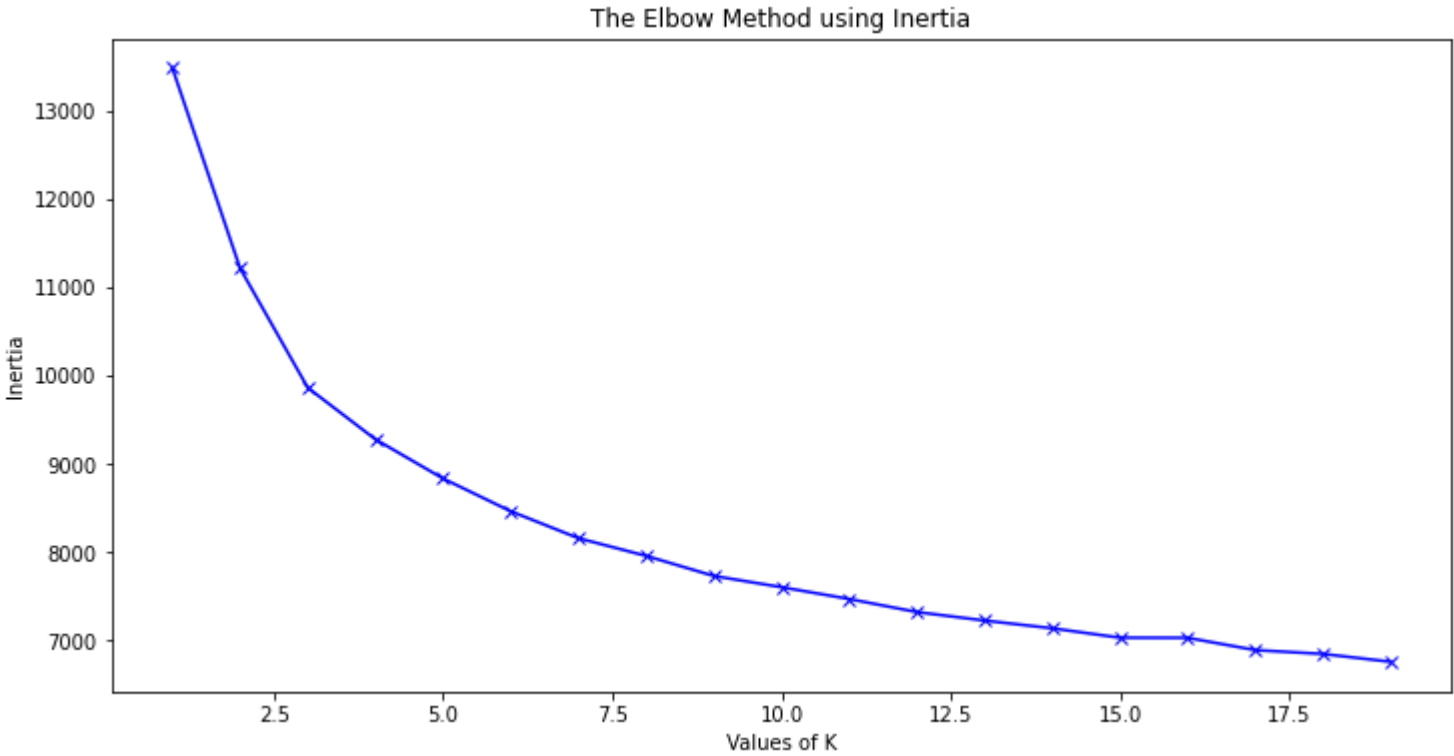
PCA

```
1 distortions = []
2 inertias = []
3 mapping3 = {}
4 mapping4 = {}
5 K2 = range(1,20)
6
7 for k in K2:
8     #Building and fitting the model
9     kmeanModel2 = KMeans(n_clusters=k).fit(df)
10    kmeanModel2.fit(df)
11
12    distortions.append(sum(np.min(cdist(df, kmeanModel2.cluster_centers_,
13                                     'euclidean'),axis=1)) / df.shape[0])
14    inertias.append(kmeanModel2.inertia_)
15
16    mapping3[k] = sum(np.min(cdist(df, kmeanModel2.cluster_centers_,
17                                'euclidean'),axis=1)) / df.shape[0]
18    mapping4[k] = kmeanModel2.inertia_
```

```
1 plt.plot(K2, distortions, 'bx-')
2 plt.xlabel('Values of K')
3 plt.ylabel('Distortion')
4 plt.title('The Elbow Method using Distortion')
5 plt.show()
```



```
1 plt.plot(K2, inertias, 'bx-')
2 plt.xlabel('Values of K')
3 plt.ylabel('Inertia')
4 plt.title('The Elbow Method using Inertia')
5 plt.show()
```

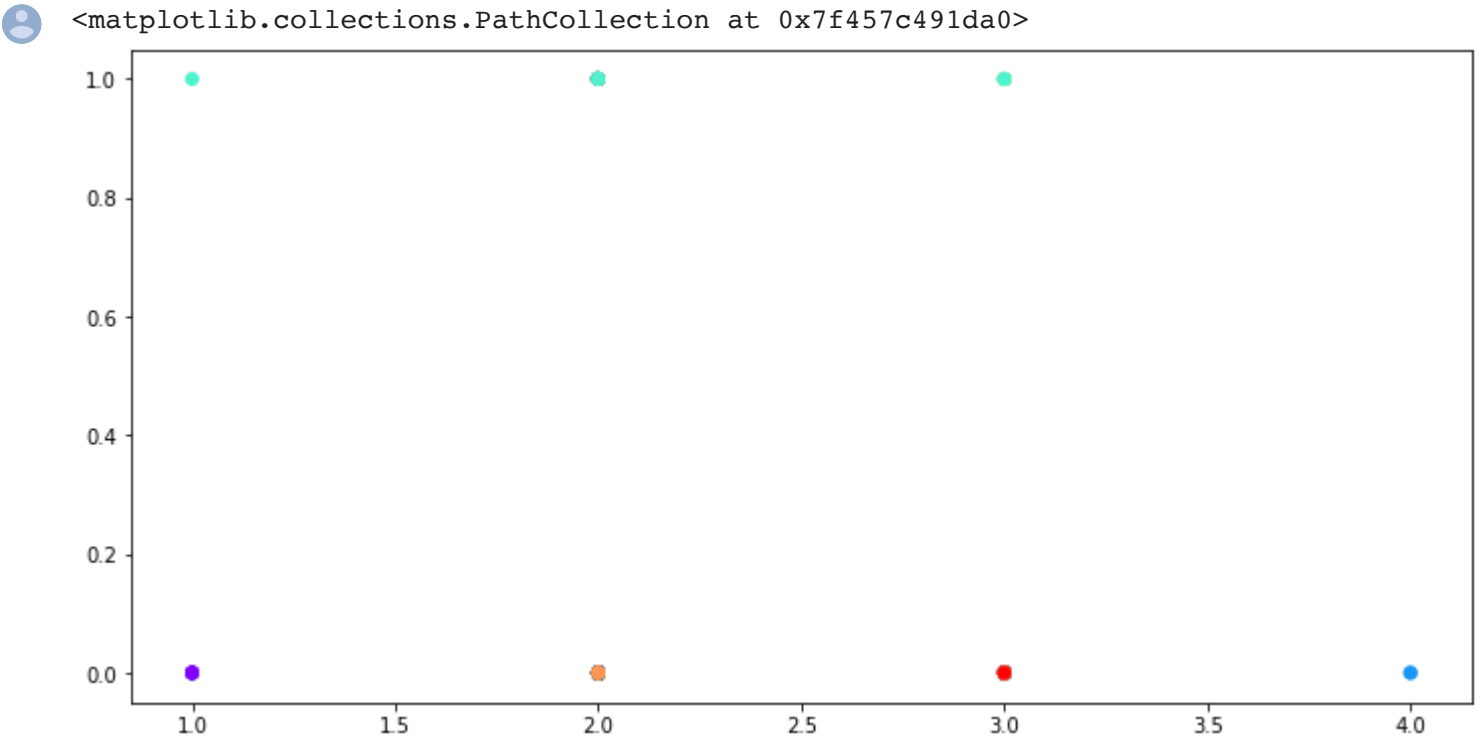


```
1 We decided to use 6 clusters.
```

```
1 #fit the kmeans model
2 kmeans2= KMeans(n_clusters=6)
3 kmeans6=kmeans2.fit_predict(df)
```

```
1 df=np.array(df)
```

```
1 #visualizing clustering
2 plt.scatter(df[:,0],df[:,1], c=kmeans6, cmap='rainbow')
```



The reason why we got imperfect plot is that there are too many dimensions here.

